Ordinal Independent Variables

Richard Williams, University of Notre Dame, https://www3.nd.edu/~rwilliam/
Last revised February 29, 2020

We often want to use ordinal variables as independent/explanatory variables in our models. Rightly or wrongly, it is very common to treat such variables as continuous. Or, more precisely, as having interval-level measurement with linear effects. When the items uses a Likert scale (e.g. Strongly Disagree, Disagree, Neutral, Agree, Strongly Agree) this may be a reasonable practice. However, many ordinal items use categories that clearly are not equally spaced, e.g. the options might be something like "daily," "a few times a week," "once a week", "a few times a month",... "once a year", "never."

In the paper referenced above, David J. Pasta makes a strong case for (usually) treating ordinal variables as continuous, even when the spacing is not equal across categories. He says (pp. 2 -3)

> One concern often expressed is that "we don't know that the ordinal categories are equally spaced." That is true enough – we don't. But we also don't "know" that the relationship between continuous variables is linear, which means we don't "know" that a one-unit change in a continuous variable has the same effect no matter whether it is a change between two relatively low values or a change between two relatively high values. In fact, when it's phrased that way -- rather than "is the relationship linear?" -- I find a lot more uncertainty in my colleagues. It turns out that it doesn't matter that much in practice – the results are remarkably insensitive to the spacing of an ordinal variable except in the most extreme cases. It does, however, matter more when you consider the products of ordinal variables.
>
> I am squarely in the camp that says "everything is linear to a first approximation" and therefore I am very cheerful about treating ordinal variables as continuous. Deviations from linearity can be important and should be considered once you have the basics of the model established, but it is very rare for an ordinal variable to be an important predictor and have it not be important when considered as a continuous variable. That would mean that the linear component of the relationship is negligible but the non-linear component is substantial. It is easy to create artificial examples of this situation, but they are very, very rare in practice.

To elaborate on one of Pasta's points – Even variables with interval-level coding don't necessarily have linear effects. You may need to take logs, add squared terms, estimate spline functions, etc. I think the issue is just a bit more obvious with ordinal variables because the number of possible values is limited and it is often questionable to believe that the categories are equally spaced.

Long and Freese (in the 2006 edition of their book) agree that ordinal variables are often treated as continuous. But they add (p. 421) that

The advantage of this approach is that interpretation is simpler, but to take advantage of this simplicity you must make the strong assumption that successive categories of the ordinal independent variable are equally spaced. For example, it implies that an increase from no publications by the mentor to a few publications involves an increase of the same amount of productivity as an increase from a few to some, from some to many, and from many to lots of publications. Accordingly, before treating an ordinal independent variable as if it were interval, you should test whether this leads to a loss of formation about the association between the independent and dependent variable.

In short, it will often be ok to treat an ordinal variable as though it had linear effects. The greater parsimony that results from doing so may be enough to offset any disadvantages that result. But, there are ways to formally test whether the assumption of linearity is justified.

Likelihood Ratio Chi-Square Test and/or BIC tests. Here, you estimate two models. In the constrained model the ordinal variable is treated as continuous, in the unconstrained model it is treated as categorical. You then use an LR chi-square test (or a BIC test or AIC test) to decide whether use of the more parsimonious continuous measure is justified.

```
. webuse nhanes2f, clear
. * Treat health as continuous
. logit diabetes c.health, nolog

Logistic regression                             Number of obs    =      10,335
                                                LR chi2(1)       =      428.14
                                                Prob > chi2      =      0.0000
Log likelihood = -1784.9973                     Pseudo R2        =      0.1071


------------------------------------------------------------------------------
    diabetes |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-------------+----------------------------------------------------------------
      health |  -.8128706   .0421577   -19.28   0.000    -.8954982    -.730243
       _cons |   -.637503   .1113932    -5.72   0.000    -.8558296   -.4191764
------------------------------------------------------------------------------

. est store m1

. * Now treat health as categorical
. logit diabetes i.health, nolog

Logistic regression                             Number of obs    =      10,335
                                                LR chi2(4)       =      429.74
                                                Prob > chi2      =      0.0000
Log likelihood = -1784.1984                     Pseudo R2        =      0.1075


------------------------------------------------------------------------------
    diabetes |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-------------+----------------------------------------------------------------
      health |
        fair |  -.7493387   .1262017    -5.94   0.000    -.9966895   -.5019878
     average |  -1.567205   .1302544   -12.03   0.000    -1.822499   -1.311911
        good |  -2.554012   .1780615   -14.34   0.000    -2.903006   -2.205018
   excellent |  -3.116457   .2262238   -13.78   0.000    -3.559848   -2.673067
             |
       _cons |  -1.481605   .0953463   -15.54   0.000     -1.66848   -1.294729
------------------------------------------------------------------------------

. est store m2
```

```
. * Now do LR/ BIC/ AIC tests
. lrtest m1 m2, stats

Likelihood-ratio test                               LR chi2(3)  =       1.60
(Assumption: m1 nested in m2)                       Prob > chi2 =     0.6599

Akaike's information criterion and Bayesian information criterion

-----------------------------------------------------------------------------
      Model |        Obs  ll(null)  ll(model)      df         AIC         BIC
------------+----------------------------------------------------------------
         m1 |     10,335 -1999.067  -1784.997       2    3573.995    3588.481
         m2 |     10,335 -1999.067  -1784.198       5    3578.397    3614.613
-----------------------------------------------------------------------------
              Note: N=Obs used in calculating BIC; see [R] BIC note.
```

A visual inspection of the coefficients from the 2nd model indeed suggests that the effects of health are continuous, i.e. each coefficient is about .75 greater than the coefficient before it. The LR/ BIC/ AIC tests also all agree that the more parsimonious model that treats health as a continuous variable is preferable.

Wald tests. Of course, you can't always do LR tests. Luckily, Wald tests are also possible. One way to do this is by including both the continuous and categorical versions of the ordinal variable in the analysis. If the effects of the categorical variable are not statistically significant, then the continuous version alone is sufficient. Note that, because we are including two versions of the ordinal variable, two categories of the ordinal variable must be excluded rather than the usual one. We can do this via use of the o. notation (o stands for omitted).

```
. * Wald test
. logit diabetes c.health o(1 2).health, nolog

Logistic regression                             Number of obs   =      10,335
                                                LR chi2(4)      =      429.74
                                                Prob > chi2     =      0.0000
Log likelihood = -1784.1984                     Pseudo R2       =      0.1075

-----------------------------------------------------------------------------
    diabetes |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-------------+---------------------------------------------------------------
      health |  -.7493387   .1262017    -5.94   0.000    -.9966895   -.5019878
             |
      health |
        fair |          0  (omitted)
     average |  -.0685278   .2104996    -0.33   0.745    -.4810995    .3440439
        good |  -.3059957   .3471345    -0.88   0.378    -.9863668    .3743754
   excellent |  -.1191024   .4829907    -0.25   0.805    -1.065747    .8275419
             |
       _cons |  -.7322659   .2078451    -3.52   0.000    -1.139635   -.3248969
-----------------------------------------------------------------------------

. testparm i.health

 ( 1)  [diabetes]3.health = 0
 ( 2)  [diabetes]4.health = 0
 ( 3)  [diabetes]5.health = 0

           chi2(  3) =      1.56
         Prob > chi2 =    0.6689
```

Again, the results indicate that the continuous version of the variable is fine.

Other options. Other strategies for dealing with ordinal independent variables have been proposed. In the Statalist thread linked to above, Ben Earnhart notes that a common, if not necessarily correct approach, is to code the midpoint of categories, e.g.

```
"daily"=1
"once a week"=1/7
"a few times a month"=(1/7)*(3/4)
"once a year"=1/365.25
"never"=0
```

Of course, some categories may be open-ended (e.g. $100,000 or more) which can make this strategy problematic.

Maarten Buis (same thread) also suggests that the `sheafcoef` command (available from SSC) can be used. The help file for `sheafcoef` says

> sheafcoef is a post-estimation command that estimates sheaf coefficients (Heise 1972). A sheaf coefficient assumes that a block of variables influence the dependent variable through a latent variable. sheafcoef displays the effect of the latent variable and the effect of the observed variables on the latent variable. The scale of the latent variable is identified by setting the standard deviation equal to one. The origin of the latent variable is identified by setting it to zero when all observed variables in its block are equal to zero. This means that the mean of the latent variable is not (necessarily) equal to zero. The final identifying assumption is that the effect of the latent variable is always positive, so to give a substantive interpretation of the direction of the effect, one needs to look at the effects of the observed variables on the latent variable. Alternatively, one can specify one "key" variable in each block of variables, which identifies the direction of a latent variable, either by specifying that the latent variable has a high value when the key variable has a high value or that the latent variable has a low value when the key variable has a high value.

> The assumption that the effect of a block of variables occurs through a latent variable is not a testable constraint; it is just a different way of presenting the results from the original model. Its main usefulness is in comparing the relative strength of the influence of several blocks of variables. For example, say we want to know what determines the probability of working non-standard hours and we have a block of variables representing characteristics of the job and another block of variables representing the family situation of the respondent, and we want to say something about the relative importance of job characteristics versus family situation. In that case one could estimate a logit model with both blocks of variables and optionally some other control variables. After that one can use sheafcoef to display the effects of two latent variables, family background and job characteristics, which are both standardized to have a standard deviation of 1, and can thus be more easily compared.

> The output is divided into a number of equations. The top equation, labeled "main", represents the effects of the latent variables and other control variables (if any) on the dependent variable. The names of the latent variables are as specified in the latent() option. If no names are specified, they will be called "lvar1", "lvar2", etc. Below the main equation, one additional equation for every latent variable is displayed, labelled "on_name1", "on_name2", etc., where "name1" and "name2" are the names of the latent variables. These are the effects of the observed variables on the latent variable.

The sheaf coeficients and the variance covariance matrix of all the coefficients are estimated using nlcom. sheafcoef can be used after any regular estimation command (that is, a command that leaves its results behind in e(b) and e(V)), The only constraint is that the observed variables that make up the latent variable(s) must all come from the same equation.

Here is an example. As far as I know, `sheafcoef` does not support factor variables, so we have to compute the dummies ourselves.

```
. * sheaf coefficients
. tab health, gen(hlth)

1=poor,..., |
5=excellent |      Freq.     Percent        Cum.
------------+-----------------------------------
      poor |        729        7.05        7.05
      fair |      1,670       16.16       23.21
   average |      2,938       28.43       51.64
      good |      2,591       25.07       76.71
 excellent |      2,407       23.29      100.00
------------+-----------------------------------
     Total |     10,335      100.00

. logit diabetes hlth2 hlth3 hlth4 hlth5, nolog

Logistic regression                             Number of obs   =      10,335
                                                LR chi2(4)      =      429.74
                                                Prob > chi2     =      0.0000
Log likelihood = -1784.1984                     Pseudo R2       =      0.1075

------------------------------------------------------------------------------
    diabetes |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-------------+----------------------------------------------------------------
       hlth2 |  -.7493387   .1262017    -5.94   0.000    -.9966895   -.5019878
       hlth3 |  -1.567205   .1302544   -12.03   0.000    -1.822499   -1.311911
       hlth4 |  -2.554012   .1780615   -14.34   0.000    -2.903006   -2.205018
       hlth5 |  -3.116457   .2262238   -13.78   0.000    -3.559848   -2.673067
       _cons |  -1.481605   .0953463   -15.54   0.000     -1.66848   -1.294729
------------------------------------------------------------------------------

. sheafcoef, latent(hlth: hlth2 hlth3 hlth4 hlth5)
------------------------------------------------------------------------------
    diabetes |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-------------+----------------------------------------------------------------
main         |
        hlth |   .9751702   .0668694    14.58   0.000     .8441085    1.106232
       _cons |  -1.481605   .0953463   -15.54   0.000     -1.66848   -1.294729
-------------+----------------------------------------------------------------
on_hlth      |
       hlth2 |  -.7684183   .1401325    -5.48   0.000    -1.043073   -.4937637
       hlth3 |  -1.607109   .1688974    -9.52   0.000    -1.938142   -1.276077
       hlth4 |  -2.619042   .1967101   -13.31   0.000    -3.004587   -2.233497
       hlth5 |  -3.195808   .1154059   -27.69   0.000       -3.422   -2.969617
------------------------------------------------------------------------------
```

In short, the main equation tells us how the underlying latent variable hlth affects the dependent variable diabetes. The on_hlth equation shows you how the observed hlth dummies affect the latent variable hlth. You don't need to assume that the categories are equally spaced.

Here is another example. In this case the LR test says we should NOT treat the ordinal variable agegrp as continuous. Visual inspection of the coefficients in the model that treats agegrp as categorical also suggests that it may not be correct to treat the effects of the variable as linear. However the BIC test disagrees, so a reasonable case could be made for going with the more parsimonious model.

```
. * Another example: agegrp
. logit diabetes c.agegrp, nolog

Logistic regression                             Number of obs   =      10,335
                                                LR chi2(1)      =      326.98
                                                Prob > chi2     =      0.0000
Log likelihood = -1835.5776                     Pseudo R2       =      0.0818

------------------------------------------------------------------------------
    diabetes |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-------------+----------------------------------------------------------------
      agegrp |   .5533155   .0350691    15.78   0.000     .4845813    .6220497
       _cons |  -5.216729   .1683673   -30.98   0.000    -5.546723   -4.886735
------------------------------------------------------------------------------

. est store m1
. logit diabetes i.agegrp, nolog

Logistic regression                             Number of obs   =      10,335
                                                LR chi2(5)      =      337.17
                                                Prob > chi2     =      0.0000
Log likelihood = -1830.4836                     Pseudo R2       =      0.0843

------------------------------------------------------------------------------
    diabetes |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-------------+----------------------------------------------------------------
      agegrp |
     age30-39 |   .7021745   .3396247     2.07   0.039     .0365223    1.367827
     age40-49 |   1.660128   .3028614     5.48   0.000      1.06653    2.253725
     age50-59 |   2.207308   .2860264     7.72   0.000     1.646706    2.767909
     age60-69 |    2.63842   .2677401     9.85   0.000     2.113659     3.16318
      age 70+ |   2.971236   .2779455    10.69   0.000     2.426472    3.515999
             |
       _cons |  -5.034786   .2590377   -19.44   0.000     -5.54249   -4.527081
------------------------------------------------------------------------------

. est store m2
. lrtest m1 m2, stats

Likelihood-ratio test                           LR chi2(4)  =      10.19
(Assumption: m1 nested in m2)                   Prob > chi2 =     0.0374

Akaike's information criterion and Bayesian information criterion

-----------------------------------------------------------------------------
       Model |        Obs  ll(null)  ll(model)     df         AIC         BIC
-------------+---------------------------------------------------------------
          m1 |     10,335 -1999.067  -1835.578      2    3675.155    3689.642
          m2 |     10,335 -1999.067  -1830.484      6    3672.967    3716.427
-----------------------------------------------------------------------------
           Note: N=Obs used in calculating BIC; see [R] BIC note.
```

Using `sheafcoef`,

```
. * Sheaf coefficients for agegrp
. tab agegrp, gen(xage)

 Age groups |
        1-6 |      Freq.     Percent        Cum.
------------+-----------------------------------
    age20-29 |      2,320       22.44       22.44
    age30-39 |      1,621       15.68       38.13
    age40-49 |      1,270       12.29       50.41
    age50-59 |      1,289       12.47       62.88
    age60-69 |      2,852       27.59       90.47
     age 70+ |        985        9.53      100.00
------------+-----------------------------------
       Total |     10,337      100.00

. logit diabetes xage2 xage3 xage4 xage5 xage6, nolog

Logistic regression                             Number of obs   =      10,335
                                                LR chi2(5)      =      337.17
                                                Prob > chi2     =      0.0000
Log likelihood = -1830.4836                     Pseudo R2       =      0.0843


------------------------------------------------------------------------------
    diabetes |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-------------+----------------------------------------------------------------
       xage2 |   .7021745   .3396247     2.07   0.039     .0365223    1.367827
       xage3 |   1.660128   .3028614     5.48   0.000      1.06653    2.253725
       xage4 |   2.207308   .2860264     7.72   0.000     1.646706    2.767909
       xage5 |    2.63842   .2677401     9.85   0.000     2.113659     3.16318
       xage6 |   2.971236   .2779455    10.69   0.000     2.426472    3.515999
       _cons |  -5.034786   .2590377   -19.44   0.000     -5.54249   -4.527081
------------------------------------------------------------------------------

. sheafcoef, latent(age: xage2 xage3 xage4 xage5 xage6)
------------------------------------------------------------------------------
    diabetes |      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-------------+----------------------------------------------------------------
main         |
         age |   1.106507   .0915181    12.09   0.000     .9271344    1.285879
       _cons |  -5.034786   .2590377   -19.44   0.000     -5.54249   -4.527081
-------------+----------------------------------------------------------------
on_age       |
       xage2 |   .6345868   .2841502     2.23   0.026     .0776627    1.191511
       xage3 |   1.500333   .1910889     7.85   0.000     1.125805     1.87486
       xage4 |   1.994844   .1405728    14.19   0.000     1.719326    2.270362
       xage5 |   2.384459   .0891692    26.74   0.000     2.209691    2.559227
       xage6 |    2.68524   .1076525    24.94   0.000     2.474245    2.896235
------------------------------------------------------------------------------
```

In short, with `sheafcoef`, we potentially get the advantages of treating an ordinal variable as continuous, without actually having to assume that categories are equally spaced. Whether it is worth the trouble is another matter; you can judge based on the circumstances. It may depend on what the tests of linear effects say or how reasonable it is to treat a variable as continuous based on its coding.